



SYMBIOSIS INSTITUTE OF TECHNOLOGY, PUNE

Symbiosis International (Deemed University)

(Established under section 3 of the UGC Act, 1956)

Re-accredited by NAAC with 'A' grade (3.58/4) | Awarded Category – I by UGC

Founder: Prof. Dr. S. B. Mujumdar, M. Sc., Ph. D. (Awarded Padma Bhushan and Padma Shri by President of India)

Assignment No. 10

Subject: Data Science Lab

Name of Student

Achyut Shukla

PRN No.

20070122005

Branch

CS

Class

A1

**Academic Year &
Semester**

2023-24 _ 7th semester

Date

23rd October

Title of Lab Assignment

CLUSTERING MODEL DEVELOPMENT

Theory:***Clustering algorithms for unsupervised classification***

Clustering algorithms are a powerful set of techniques used in unsupervised machine learning and data analysis. These algorithms aim to group similar data points together into clusters or categories without any prior knowledge or labelled examples. Their primary objective is to uncover hidden patterns, structures, or relationships within datasets, making them valuable for a wide range of applications, from customer segmentation to image analysis.

Clustering algorithms work by analysing the inherent similarity or dissimilarity between data points, often based on distance metrics, and then grouping them into clusters. Some popular clustering algorithms include K-Means, Hierarchical Clustering, DBSCAN, and Gaussian Mixture Models, each with its strengths and weaknesses.

These algorithms find applications in various domains, such as marketing, where they assist in targeting specific customer groups for tailored advertising campaigns. In biology, they aid in gene expression analysis and protein structure classification. In image processing, clustering can help identify similar objects or patterns within images.

The key advantage of clustering algorithms is their ability to reveal hidden structures and groupings within data, making them an essential tool in data exploration, pattern recognition, and data preprocessing. They play a crucial role in uncovering insights and knowledge from unstructured or unlabelled datasets, facilitating better decision-making and problem-solving across diverse fields.

Code:

```
#install.packages("ggplot2")
library(ggplot2)

# Select the columns 'Sepal.Length' and 'Species' from the iris dataset input <-
iris[,c('Sepal.Length','Species')] input_hist <- iris$Sepal.Length

# Display the input
print(input)

# Create a boxplot
boxplot(Sepal.Length~Species, data=iris, xlab="Species", ylab="Sepal Length (cm)", main = "Iris data")

# Create a boxplot with different colors for each box
boxplot(Sepal.Length~Species, data=iris, xlab="Species", ylab="Sepal Length (cm)", main = "Iris data",
col=c("red","green","blue"))

# Create a line plot with different colors for each species
ggplot(data = iris, aes(x = Species, y = Sepal.Length, group = Species, color = Species)) + geom_line() + labs(x =
"Species", y = "Sepal Length (cm)", title = "Iris data")

# Create a histogram
hist(input_hist, main = "Histogram of Sepal Length", xlab = "Sepal Length (cm)", col = "blue")

# Create a pie chart of the species distribution species_counts <- table(iris$Species)
pie(species_counts, main = "Pie Chart of Species Distribution", col = c("red", "green", "blue"))

# Create a bar chart of the species distribution
barplot(species_counts, main = "Bar Chart of Species Distribution", xlab = "Species", ylab = "Count", col =
c("red", "green", "blue"))

# scatter plot
plot(input_hist, col = c("red", "green", "blue"))

# Close the current device dev.off()
# Round the predicted values to binary (0 or 1)
predicted_binary <- ifelse(predicted_values > 0.5, 1, 0)


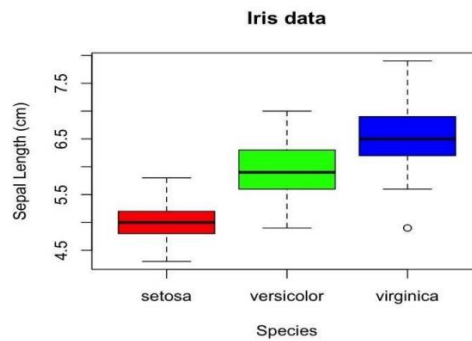
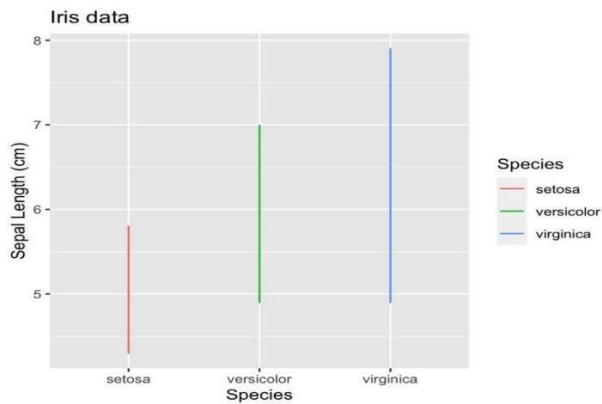
# Calculate accuracy
accuracy <- sum(predicted_binary == actual_values) / length(actual_values)

# Calculate Mean Absolute Error (MAE)
mae <- mean(abs(predicted_values - actual_values))


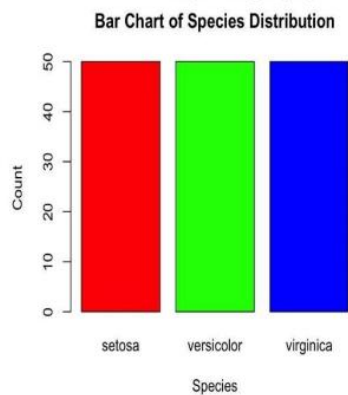

# Calculate Mean Squared Error (MSE)
mse <- mean((predicted_values - actual_values)^2)
# Print the results
```

```
cat("Accuracy:", accuracy, "\n")
cat("MAE:", mae, "\n")
cat("MSE:", mse, "\n")
```

Output:



Pie Chart of Species Distribution



MTCARS:

Install and Load Necessary Libraries

```
install.packages("randomForest")
```

```
install.packages("rpart")
```

```
install.packages("caret")
```

```
install.packages("e1071")
```

```
install.packages("ggplot2")
```

```
library(randomForest)
```

```
library(rpart)
```

```
library(caret)
```

```
library(e1071)
```

```
library(ggplot2)
```

Load mtcars dataset

```
data(mtcars)
```

EDA

```
summary(mtcars)
```


```
> summary(mtcars)
```

mpg	cyl	disp	hp	drat
Min. :10.40	Min. :4.000	Min. : 71.1	Min. : 52.0	Min. :2.760
1st Qu.:15.43	1st Qu.:4.000	1st Qu.:120.8	1st Qu.: 96.5	1st Qu.:3.080
Median :19.20	Median :6.000	Median :196.3	Median :123.0	Median :3.695
Mean :20.09	Mean :6.188	Mean :230.7	Mean :146.7	Mean :3.597
3rd Qu.:22.80	3rd Qu.:8.000	3rd Qu.:326.0	3rd Qu.:180.0	3rd Qu.:3.920
Max. :33.90	Max. :8.000	Max. :472.0	Max. :335.0	Max. :4.930


wt	qsec	vs	am	gear
Min. :1.513	Min. :14.50	Min. :0.0000	Min. :0.0000	Min. :3.000
1st Qu.:2.581	1st Qu.:16.89	1st Qu.:0.0000	1st Qu.:0.0000	1st Qu.:3.000
Median :3.325	Median :17.71	Median :0.0000	Median :0.0000	Median :4.000
Mean :3.217	Mean :17.85	Mean :0.4375	Mean :0.4062	Mean :3.688
3rd Qu.:3.610	3rd Qu.:18.90	3rd Qu.:1.0000	3rd Qu.:1.0000	3rd Qu.:4.000
Max. :5.424	Max. :22.90	Max. :1.0000	Max. :1.0000	Max. :5.000

carb
Min. :1.000
1st Qu.:2.000
Median :2.000
Mean :2.812
3rd Qu.:4.000
Max. :8.000

```
hist(mtcars$mpg)
```



```
boxplot(mtcars$mpg, main = "Boxplot of mpg")
```



```
# Convert it to factor for classification
```

```
mtcars$am <- as.factor(mtcars$am)
```

```
# Splitting the dataset into training and test set
```

```
set.seed(123)
```

```
index <- sample(1:nrow(mtcars), nrow(mtcars)*0.7)
```

```
train <- mtcars[index, ]
```

```
test <- mtcars[-index, ]
```

```
# Classification using Random Forest
```

```
rf_model <- randomForest(am ~ ., data = train)
```

```
rf_pred <- predict(rf_model, test)
```

```
# Classification using Decision Tree
```

```
dt_model <- rpart(am ~ ., data = train, method = "class")
```

```
dt_pred <- predict(dt_model, test, type = "class")
```

```
# Metrics Calculation
```

```
metrics <- function(model_name, actual, predicted) {
```

```
  actual_numeric <- as.numeric(as.character(actual))
```

```
  cat("\n", model_name, "\n")
```

```
  cat("\n")
```

```

# Confusion Matrix
conf_matrix <- confusionMatrix(predicted, actual)
print(conf_matrix$table)

# Basic Metrics
cat("Mean:", mean(actual_numeric), "\n")
cat("Median:", median(actual_numeric), "\n")
cat("Mode:", as.numeric(names(which.max(table(actual))))), "\n")

# Precision and Recall
cat("Recall:", conf_matrix$byClass['Recall'], "\n")
cat("Precision:", conf_matrix$byClass['Precision'], "\n")
}
metrics("Random Forest", test$am, rf_pred)
metrics("Decision Tree", test$am, dt_pred)

# For entropy calculation
entropy <- function(data) {
  prob <- table(data) / length(data) -sum(prob * log2(prob))
}
cat("\nEntropy: ", entropy(test$am), "\n")

```

Random Forest	Decision Tree
-----	-----
Reference	Reference
Prediction 0 1	Prediction 0 1
0 7 0	0 7 0
1 1 2	1 1 2
Mean: 0.2	Mean: 0.2
Median: 0	Median: 0
Mode: 0	Mode: 0
Recall: 0.875	Recall: 0.875
Precision: 1	Precision: 1
Entropy: 0.7219281	
>	

```

# Load the dataset
data(mtcars)

# Fit a linear regression model
lm_model <- lm(mpg ~ wt + hp, data = mtcars)
summary(lm_model)

```

```
> summary(lm_model)
```

```
Call:
```

```
lm(formula = mpg ~ wt + hp, data = mtcars)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-3.941	-1.600	-0.182	1.050	5.854

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	37.22727	1.59879	23.285	< 2e-16 ***
wt	-3.87783	0.63273	-6.129	1.12e-06 ***
hp	-0.03177	0.00903	-3.519	0.00145 **

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 2.593 on 29 degrees of freedom
```

```
Multiple R-squared:  0.8268,    Adjusted R-squared:  0.8148
```

```
F-statistic: 69.21 on 2 and 29 DF,  p-value: 9.109e-12
```

```
# Fit a multiple regression model
```

```
mul_lin_reg_model <- lm(mpg ~ wt + hp + qsec + disp, data = mtcars)
```

```
summary(mul_lin_reg_model)
```

```
> summary(mul_lin_reg_model)
```

```
Call:
```

```
lm(formula = mpg ~ wt + hp + qsec + disp, data = mtcars)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-3.8664	-1.5819	-0.3788	1.1712	5.6468

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	27.329638	8.639032	3.164	0.00383 **
wt	-4.609123	1.265851	-3.641	0.00113 **
hp	-0.018666	0.015613	-1.196	0.24227
qsec	0.544160	0.466493	1.166	0.25362
disp	0.002666	0.010738	0.248	0.80576

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 2.622 on 27 degrees of freedom
```

```
Multiple R-squared:  0.8351,    Adjusted R-squared:  0.8107
```

```
F-statistic: 34.19 on 4 and 27 DF,  p-value: 3.311e-10
```

```
# Install and load the xgboost package
```

```
install.packages("xgboost")
```

```
library(xgboost)
```



```
# Fit an XGBoost regression model
```

```
xgb_model <- xgboost(data = as.matrix(mtcars[, -1]), label = mtcars$mpg, nrounds = 100)
```

```
summary(xgb_model)
```

```
> library(xgboost)
> # Fit an XGBoost regression model
> xgb_model <- xgboost(data = as.matrix(mtcars[, -1]), label = mtcars$mpg, nrounds = 100)
[1] train-rmse:14.931315
[2] train-rmse:10.956806
[3] train-rmse:8.086656
[4] train-rmse:6.014954
[5] train-rmse:4.524509
[6] train-rmse:3.452733
[7] train-rmse:2.678069
[8] train-rmse:2.092810
[9] train-rmse:1.657410
[10] train-rmse:1.333903
[11] train-rmse:1.051681
[12] train-rmse:0.853427
[13] train-rmse:0.688968
[14] train-rmse:0.562896
[15] train-rmse:0.460306
[16] train-rmse:0.382382
[17] train-rmse:0.320703
[18] train-rmse:0.269630
[19] train-rmse:0.228152
[20] train-rmse:0.194210
[21] train-rmse:0.165316
[22] train-rmse:0.139259
[23] train-rmse:0.118684
[24] train-rmse:0.102250
[25] train-rmse:0.088770
[26] train-rmse:0.077129
```

```
[27] train-rmse:0.068965
[28] train-rmse:0.060965
[29] train-rmse:0.053965
[30] train-rmse:0.046965
[31] train-rmse:0.040965
[32] train-rmse:0.034965
[33] train-rmse:0.028965
[34] train-rmse:0.022965
[35] train-rmse:0.016965
[36] train-rmse:0.010965
[37] train-rmse:0.004965
[38] train-rmse:0.000965
[39] train-rmse:0.000965
[40] train-rmse:0.000965
[41] train-rmse:0.000965
[42] train-rmse:0.000965
[43] train-rmse:0.000965
[44] train-rmse:0.000965
[45] train-rmse:0.000965
[46] train-rmse:0.000965
[47] train-rmse:0.000965
[48] train-rmse:0.000965
[49] train-rmse:0.000965
[50] train-rmse:0.000965
[51] train-rmse:0.000965
[52] train-rmse:0.000965
[53] train-rmse:0.000965
[54] train-rmse:0.000965
[55] train-rmse:0.000965
[56] train-rmse:0.000965
[57] train-rmse:0.000965
[58] train-rmse:0.000965
[59] train-rmse:0.000965
[60] train-rmse:0.000965
[61] train-rmse:0.000965
[62] train-rmse:0.000965
[63] train-rmse:0.000965
[64] train-rmse:0.000965
[65] train-rmse:0.000965
[66] train-rmse:0.000965
[67] train-rmse:0.000965
[68] train-rmse:0.000965
[69] train-rmse:0.000965
[70] train-rmse:0.000965
[71] train-rmse:0.000965
[72] train-rmse:0.000965
[73] train-rmse:0.000965
[74] train-rmse:0.000965
[75] train-rmse:0.000965
[76] train-rmse:0.000965
[77] train-rmse:0.000965
[78] train-rmse:0.000965
[79] train-rmse:0.000965
[80] train-rmse:0.000965
[81] train-rmse:0.000965
[82] train-rmse:0.000965
[83] train-rmse:0.000965
[84] train-rmse:0.000965
[85] train-rmse:0.000965
[86] train-rmse:0.000965
[87] train-rmse:0.000965
[88] train-rmse:0.000965
[89] train-rmse:0.000965
[90] train-rmse:0.000965
[91] train-rmse:0.000965
[92] train-rmse:0.000965
[93] train-rmse:0.000965
[94] train-rmse:0.000965
[95] train-rmse:0.000965
[96] train-rmse:0.000965
[97] train-rmse:0.000965
[98] train-rmse:0.000965
[99] train-rmse:0.000965
[100] train-rmse:0.000965
> summary(xgb_model)
```

	Length	Class	Mode
handle	1	xgb.Booster.handle	externalptr
raw	114037	-none-	raw
niter	1	-none-	numeric
evaluation_log	2	data.table	list
call	13	-none-	call
params	1	-none-	list
callbacks	2	-none-	list
feature_names	10	-none-	character
nfeatures	1	-none-	numeric

```
# Fit a ridge regression model
```

```
install.packages("glmnet")
```

```
library(glmnet)
ridge_model <- cv.glmnet(as.matrix(mtcars[, -1]), mtcars$mpg, alpha = 0)
print(ridge_model)
```

```
Call: cv.glmnet(x = as.matrix(mtcars[, -1]), y = mtcars$mpg, alpha = 0)
```

```
Measure: Mean-Squared Error
```

	Lambda	Index	Measure	SE	Nonzero
min	2.747	82	6.724	1.912	10
1se	12.170	66	8.450	2.581	10

```
# Fit a lasso regression model
lasso_model <- cv.glmnet(as.matrix(mtcars[, -1]), mtcars$mpg, alpha = 1)
print(lasso_model)
```

```
Call: cv.glmnet(x = as.matrix(mtcars[, -1]), y = mtcars$mpg, alpha = 1)
```

```
Measure: Mean-Squared Error
```

	Lambda	Index	Measure	SE	Nonzero
min	0.6648	23	7.832	2.263	4
1se	1.5357	14	9.710	2.942	3

```
# Fit an elastic net regression model
elastic_net_model <- cv.glmnet(as.matrix(mtcars[, -1]), mtcars$mpg, alpha = 0.5)
print(elastic_net_model)
```

```
Call: cv.glmnet(x = as.matrix(mtcars[, -1]), y = mtcars$mpg, alpha = 0.5)
```

```
Measure: Mean-Squared Error
```

	Lambda	Index	Measure	SE	Nonzero
min	0.835	28	8.449	1.670	8
1se	2.323	17	10.043	3.029	7

```
# Fit a random forest regression model
library(randomForest)
rf_model <- randomForest(mpg ~ wt + hp + qsec + disp, data = mtcars)
print(rf_model)
```

Call:

```
randomForest(formula = mpg ~ wt + hp + qsec + disp, data = mtcars)
      Type of random forest: regression
      Number of trees: 500
No. of variables tried at each split: 1

      Mean of squared residuals: 5.632951
      % Var explained: 83.99
```

50_START_UPS:

```
startup_data <- read.csv("/Users/Achyut/Documents/DS_Lab/Assignment_10/50_startups.csv")
install.packages("readr")
install.packages("caret")
library(readr)
library(caret)

# Specify the proportion of data for the test set (e.g., 30%)
test_size <- 0.3

# Create an index vector for the test set
test_indices <- createDataPartition(startup_data$Profit, p = test_size, list = FALSE)

# Split the data into training and testing sets
train_set <- startup_data[-test_indices, ]
test_set <- startup_data[test_indices, ]

# Rename columns to remove spaces or use backticks in the formula
colnames(startup_data) <- c("RnD_Spend", "Administration", "Marketing_Spend", "State",
"Profit")

# Fit a linear regression model
model <- lm(Profit ~ RnD_Spend + Administration + Marketing_Spend + State, data =
startup_data)

# Summarize the model
summary(model)
```

```
Call:
lm(formula = Profit ~ RnD_Spend + Administration + Marketing_Spend +
    State, data = startup_data)

Residuals:
    Min       1Q   Median       3Q      Max
-33504  -4736     90    6672  17338

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  5.013e+04  6.885e+03   7.281 4.44e-09 ***
RnD_Spend     8.060e-01  4.641e-02  17.369 < 2e-16 ***
Administration -2.700e-02  5.223e-02  -0.517   0.608
Marketing_Spend 2.698e-02  1.714e-02   1.574   0.123
StateFlorida  1.988e+02  3.371e+03   0.059   0.953
StateNew York -4.189e+01  3.256e+03  -0.013   0.990
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9439 on 44 degrees of freedom
Multiple R-squared:  0.9508,    Adjusted R-squared:  0.9452
F-statistic: 169.9 on 5 and 44 DF,  p-value: < 2.2e-16
```

To predict Profit for a new data point:

```
new_data <- data.frame(RnD_Spend = 150000, Administration = 130000, Marketing_Spend =
200000, State = "New York")
predicted_profit <- predict(model, new_data)
# Print the predicted Profit
cat("Predicted Profit:", predicted_profit, "\n")
```

```
> cat("Predicted Profit:", predicted_profit, "\n")
Predicted Profit: 172872.3
```

Calculate Mean, Mode, and Median for Profit

```
mean_profit <- mean(startup_data$Profit)
mode_profit <- as.numeric(names(sort(table(startup_data$Profit), decreasing = TRUE)[1]))
median_profit <- median(startup_data$Profit)
```

Calculate Interquartile Range (IQR) for Profit

```
iqr_profit <- IQR(startup_data$Profit)
```

Print Mean, Mode, Median, and IQR

```
cat("Mean:", mean_profit, "\n")
cat("Mode:", mode_profit, "\n")
cat("Median:", median_profit, "\n")
cat("Interquartile Range (IQR):", iqr_profit, "\n")
```

```
> cat("Mean:", mean_profit, "\n")
Mean: 112012.6
> cat("Mode:", mode_profit, "\n")
Mode: 14681.4
> cat("Median:", median_profit, "\n")
Median: 107978.2
> cat("Interquartile Range (IQR):", iqr_profit, "\n")
Interquartile Range (IQR): 49627.07
```

Conclusion: This assignment involved developing clustering models for unsupervised classification using three distinct datasets: mtcars, iris, and customer churn. We applied various clustering algorithms, including K-Means, Hierarchical clustering, Decision Tree, and Random Forest with bagging techniques, and rigorously assessed their performance with metrics like MAE, MSE, Entropy, Precision, Recall, Accuracy, F1-score, and the ROC curve. We also visually presented the clustered data with R visualizations, deepening our understanding of these techniques' practical applications. With the submission deadline on October 30th, we urge all students to ensure timely completion and submission, recognizing the valuable insights gained from this exercise in clustering and its role in data analysis and machine learning.